

EXP.NO:03

DATE:

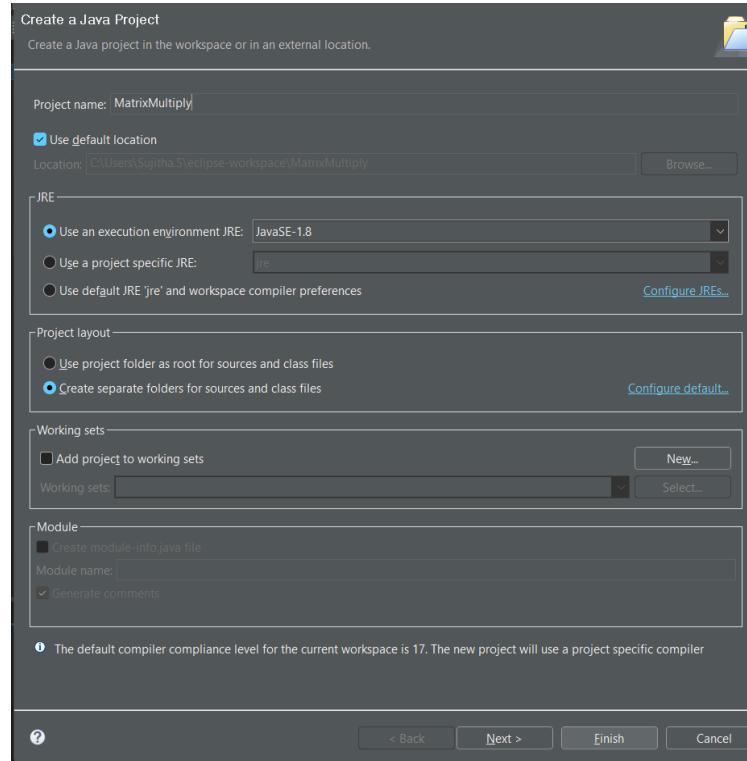
IMPLEMENTATION OF MATRIX MULTIPLICATION WITH HADOOP MAP REDUCE

AIM: To implement of Matrix Multiplication with Hadoop Map Reduce.

STEPS:

STEP 1: Run Eclipse for Java Developers

STEP 2: Create a new Java Project with name “MatrixMultiply”



STEP 3: Set the Java Environment Version to your current version of Java (JRE : 1.8)

STEP 4: Add a Package with name “com.MapReduce.java” and Create three Classes in it.

STEP 5: Create a New Class With name Map.java.

STEP 6: Now write the below program in the “Map.java” Class

PROGRAM:

```
package com.MapReduce.wc;

import org.apache.hadoop.conf.*;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
//import org.apache.hadoop.mapreduce.Mapper;
import java.io.IOException;

public class Map extends org.apache.hadoop.mapreduce.Mapper<LongWritable, Text, Text, Text>

{

    @Override

    public void map(LongWritable key, Text value, Context context)

    throws IOException, InterruptedException {

        Configuration conf = context.getConfiguration();

        int m = Integer.parseInt(conf.get("m"));

        int p = Integer.parseInt(conf.get("p"));

        String line = value.toString();

        // (M, i, j, Mij);

        String[] indicesAndValue = line.split(",");

        Text outputKey = new Text();

        Text outputValue = new Text();

        if (indicesAndValue[0].equals("M")) {

            for (int k = 0; k < p; k++) {

                outputKey.set(indicesAndValue[1] + "," + k);
```

```

// outputKey.set(i,k);

outputValue.set(indicesAndValue[0] + "," + indicesAndValue[2]
+ "," + indicesAndValue[3]);

// outputValue.set(M,j,Mij);

context.write(outputKey, outputValue);

}

} else {

// (N, j, k, Njk);

for (int i = 0; i < m; i++) {

outputKey.set(i + "," + indicesAndValue[2]); outputValue.set("N," + indicesAndValue[1] + ","
+ indicesAndValue[3]); context.write(outputKey, outputValue);

}

}

}

}

}

```

STEP 7: Now Create another class with name “Reduce.java” and paste the below program in it.

PROGRAM:

```

package com.MapReduce.wc;

import org.apache.hadoop.io.Text;

// import org.apache.hadoop.mapreduce.Reducer;

import java.io.IOException;

import java.util.HashMap;

public class Reduce

```

```

extends org.apache.hadoop.mapreduce.Reducer<Text, Text, Text, Text> { @Override

public void reduce(Text key, Iterable<Text> values, Context context)
throws IOException, InterruptedException {

String[] value;
//key=(i,k),
//Values = [(M/N,j,V/W),..]

HashMap<Integer, Float> hashA = new HashMap<Integer, Float>(); HashMap<Integer, Float> hashB =
new HashMap<Integer, Float>(); for (Text val : values) {

value = val.toString().split(",");
if (value[0].equals("M")) {

hashA.put(Integer.parseInt(value[1]), Float.parseFloat(value[2])); } else {

hashB.put(Integer.parseInt(value[1]), Float.parseFloat(value[2]));

}
}

int n = Integer.parseInt(context.getConfiguration().get("n"));

float result = 0.0f;

float m_ij;
float n_jk;

for (int j = 0; j < n; j++) {

m_ij = hashA.containsKey(j) ? hashA.get(j) : 0.0f; n_jk = hashB.containsKey(j) ? hashB.get(j) : 0.0f;
result += m_ij * n_jk;

}

if (result != 0.0f) {

context.write(null,

```

```

new Text(key.toString() + "," + Float.toString(result)));
}

}

}

```

STEP 8: Now, Create another class with name “MatrixMultiply.java” and paste the below program in it.

PROGRAM:

```

package com.MapReduce.wc;

import org.apache.hadoop.conf.*;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.*;
import org.apache.hadoop.mapreduce.*;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.input.TextInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.mapreduce.lib.output.TextOutputFormat;

public class MatrixMultiply {

    public static void main(String[] args) throws Exception {
        if (args.length != 2) {
            System.err.println("Usage: MatrixMultiply <in_dir> <out_dir>");
            System.exit(2);
        }

        Configuration conf = new Configuration();
        conf.set("m", "1000");
        conf.set("n", "100");
    }
}

```

```
conf.set("p", "1000");

@SuppressWarnings("deprecation")

Job job = new Job(conf, "MatrixMultiply");

job.setJarByClass(MatrixMultiply.class);

job.setOutputKeyClass(Text.class);

job.setOutputValueClass(Text.class);

job.setMapperClass(Map.class);

job.setReducerClass(Reduce.class);

job.setInputFormatClass(TextInputFormat.class);

job.setOutputFormatClass(TextOutputFormat.class);

FileInputFormat.addInputPath(job, new Path(args[0]));

FileOutputFormat.setOutputPath(job, new Path(args[1]));

job.waitForCompletion(true);

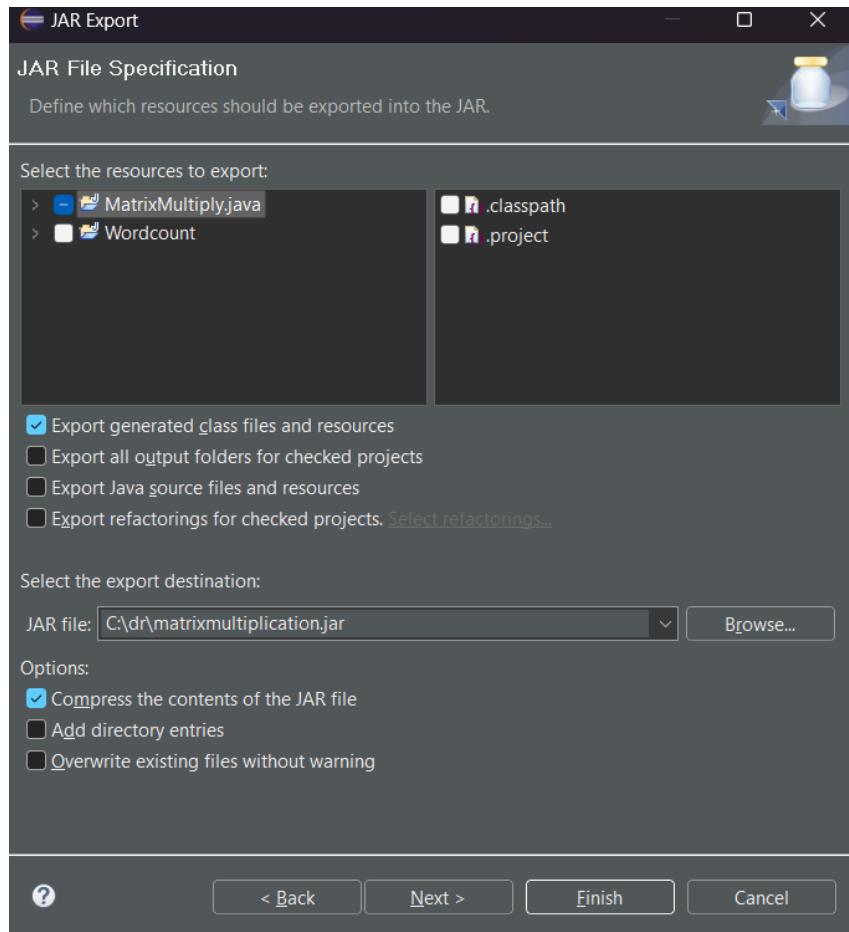
}

}
```

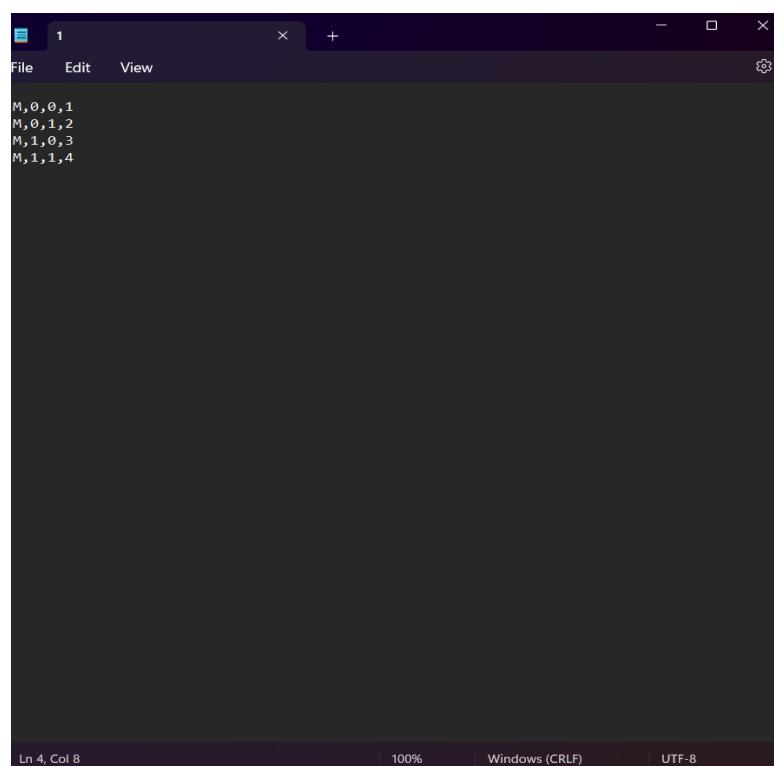
STEP 9: To resolve the errors in the programs we should add two External jar files to it.

- Hadoop_common :2.7.3.jar
- Hadoop_mapreduce:client:core:2.7.1.jar

STEP 10: Now export the project into a Jar file and name it as “MatrixMultiply.jar”



STEP 11: Now create a Text file in Notepad and name it as “1.txt” and “2.txt. write some content inside the text file and save it.



A screenshot of a dark-themed code editor window. The title bar shows the number '2' indicating two tabs. The menu bar includes 'File', 'Edit', 'View', and a settings gear icon. The main editor area contains the following text:

```
N,0,0,5
N,0,1,6
N,1,0,7
N,1,1,8|
```

The status bar at the bottom shows 'Ln 4, Col 8' on the left, zoom level '100%' in the center, and file encoding 'Windows (CRLF)' and character set 'UTF-8' on the right.

STEP 12: Now run all the deamons in Hadoop.

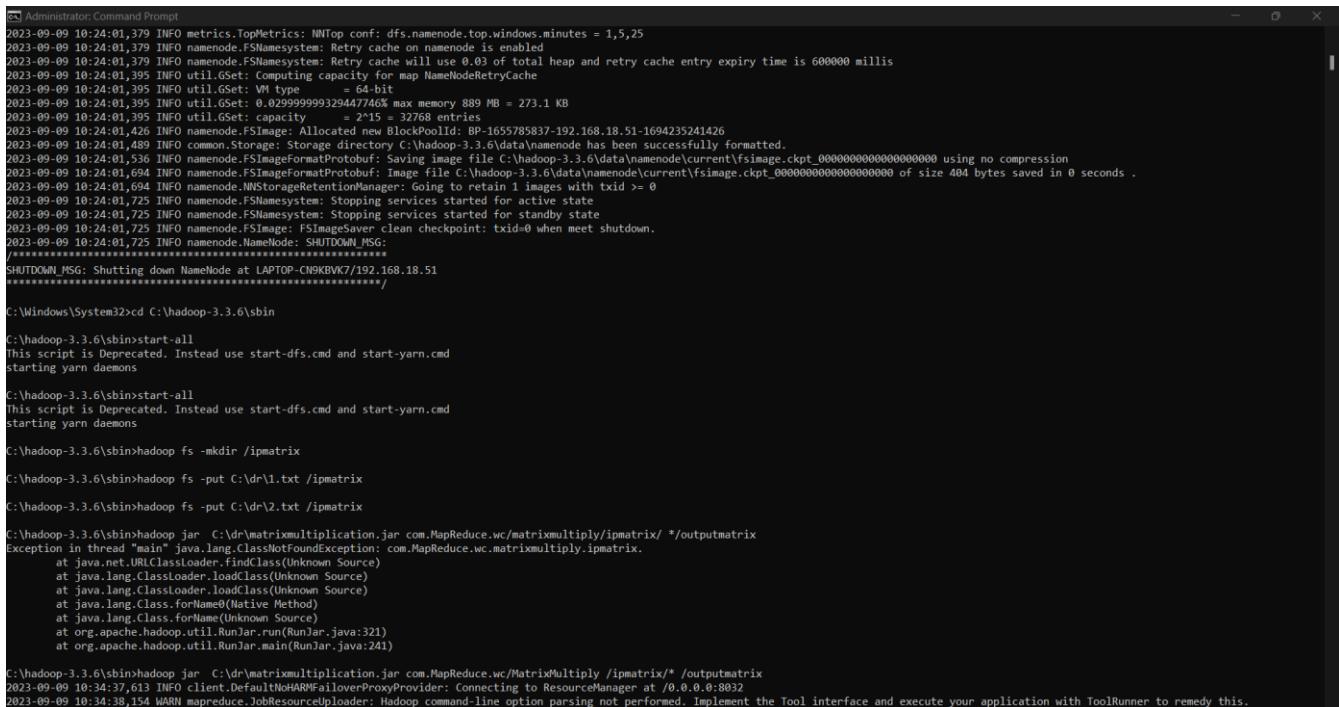
STEP 13: Create a new input directory named as “ipmatrix”.

By using the command: `hadoop fs -mkdir /ipmatrix`

STEP 14: Now put the “1.txt” and 2.txt file to the ipmatrix directory.

By using these commands: hadoop fs -put C:\dr\1.txt /ipmatrix

hadoop fs -put C:\dr\2.txt /ipmatrix

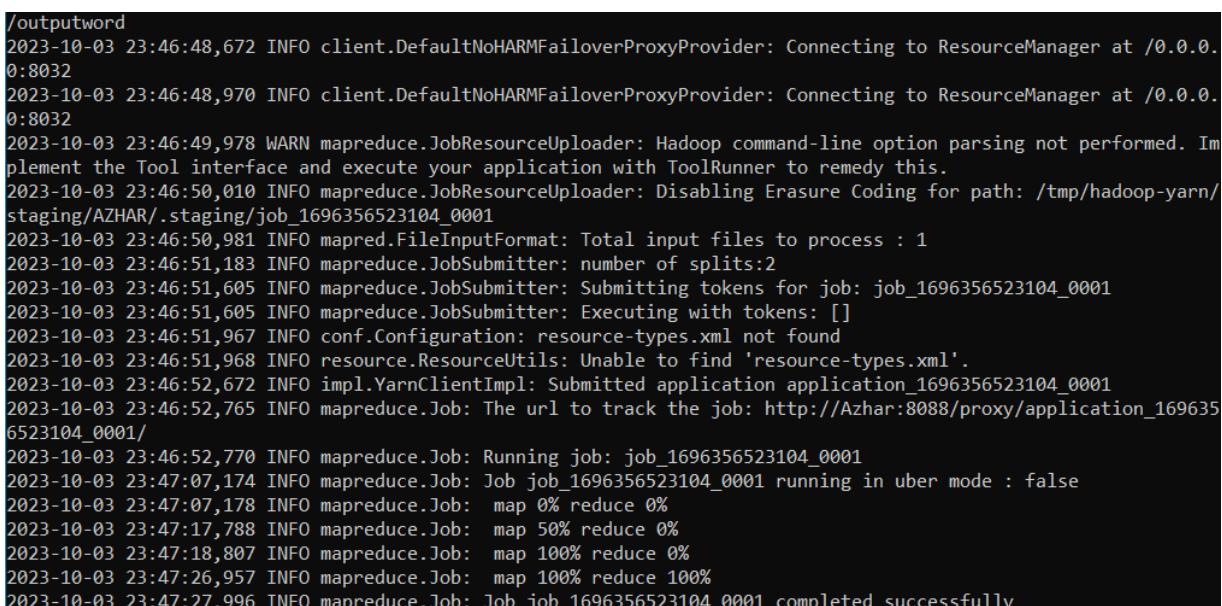


```
Administrator: Command Prompt
2023-09-09 10:24:01,379 INFO metrics.TopMetrics: NWTop conf: dfs.namenode.top.windows.minutes = 1,5,25
2023-09-09 10:24:01,379 INFO namenode.FSNamesystem: Retry cache on namenode is enabled
2023-09-09 10:24:01,379 INFO namenode.FSNamesystem: Retry cache will use 0.03 of total heap and retry cache entry expiry time is 600000 millis
2023-09-09 10:24:01,395 INFO util.GSet: Compressing capacity for map NameNodeRetryCache = 64-bit
2023-09-09 10:24:01,395 INFO util.GSet: capacity = 2^15 = 32768 entries
2023-09-09 10:24:01,426 INFO namenode.FSImage: Allocated new BlockPoolId: BP-1655785837-192.168.18.51-1694235241426
2023-09-09 10:24:01,489 INFO common.Storage: Storage directory C:\hadoop-3.3.6\data\namenode\current\fsimage has been successfully Formatted.
2023-09-09 10:24:01,528 INFO namenode.FSImageFormatProtocol: Saving image file C:\hadoop-3.3.6\data\namenode\current\fsimage.ckpt_00000000000000000000 of size 404 bytes saved in 0 seconds .
2023-09-09 10:24:01,694 INFO namenode.NMStorageRetentionManager: Going to retain 1 images with txid >= 0
2023-09-09 10:24:01,725 INFO namenode.FSNamesystem: Stopping services started for active state
2023-09-09 10:24:01,725 INFO namenode.FSNamesystem: Stopping services started for standby state
2023-09-09 10:24:01,725 INFO namenode.FSImage: FSImageSaver clean checkpoint: txid=0 when meet shutdown.
2023-09-09 10:24:01,725 INFO namenode.NameNode: SHUTDOWN_MSG:
*****Shutdown Message*****:
SHUTDOWN_MSG: Shutting down NameNode at LAPTOP-CM9KBVK7/192.168.18.51
*****
C:\Windows\System32>cd C:\hadoop-3.3.6\sbin
C:\hadoop-3.3.6\bin>start-all
This script is Deprecated. Instead use start-dfs.cmd and start-yarn.cmd
starting yarn daemons
C:\hadoop-3.3.6\bin>start-all
This script is Deprecated. Instead use start-dfs.cmd and start-yarn.cmd
starting yarn daemons
C:\hadoop-3.3.6\bin>hadoop fs -mkdir /ipmatrix
C:\hadoop-3.3.6\bin>hadoop fs -put C:\dr\1.txt /ipmatrix
C:\hadoop-3.3.6\bin>hadoop fs -put C:\dr\2.txt /ipmatrix
C:\hadoop-3.3.6\bin>Hadoop jar C:\dr\matrixmultiplication.jar com.MapReduce.wc/matrixmultiply/ipmatrix/*/outputmatrix
Exception in thread "main" java.lang.ClassNotFoundException: com.MapReduce.wc.matrixmultiply.ipmatrix.
    at java.net.URLClassLoader.findClass(Unknown Source)
    at java.lang.ClassLoader.loadClass(Unknown Source)
    at java.lang.ClassLoader.loadClass(Unknown Source)
    at java.lang.Class.forName0(Native Method)
    at java.lang.Class.forName(Unknown Source)
    at org.apache.hadoop.util.RunJar.run(RunJar.java:321)
    at org.apache.hadoop.util.RunJar.main(RunJar.java:241)

C:\hadoop-3.3.6\bin>hadoop jar C:\dr\matrixmultiplication.jar com.MapReduce.wc/MatrixMultiply /ipmatrix/* /outputmatrix
2023-09-09 10:34:37,613 INFO client.DefaultNoHARMFailoverProxyProvider: Connecting to ResourceManager at /0.0.0:8032
2023-09-09 10:34:38,154 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
```

STEP 15: Run the Jar file created from the project

Using the command: hadoop jar C:\dr\matrixmultiplication.jar com.MapReduce.wc/MatrixMultiply /ipmatrix/* /outputmatrix



```
/outputword
2023-10-03 23:46:48,672 INFO client.DefaultNoHARMFailoverProxyProvider: Connecting to ResourceManager at /0.0.0.0:8032
2023-10-03 23:46:48,970 INFO client.DefaultNoHARMFailoverProxyProvider: Connecting to ResourceManager at /0.0.0.0:8032
2023-10-03 23:46:49,978 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
2023-10-03 23:46:50,010 INFO mapreduce.JobResourceUploader: Disabling Erasure Coding for path: /tmp/hadoop-yarn/staging/AZHAR/.staging/job_1696356523104_0001
2023-10-03 23:46:50,981 INFO mapred.FileInputFormat: Total input files to process : 1
2023-10-03 23:46:51,183 INFO mapreduce.JobSubmitter: number of splits:2
2023-10-03 23:46:51,605 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1696356523104_0001
2023-10-03 23:46:51,605 INFO mapreduce.JobSubmitter: Executing with tokens: []
2023-10-03 23:46:51,967 INFO conf.Configuration: resource-types.xml not found
2023-10-03 23:46:51,968 INFO resource.ResourceUtils: Unable to find 'resource-types.xml'.
2023-10-03 23:46:52,672 INFO impl.YarnClientImpl: Submitted application application_1696356523104_0001
2023-10-03 23:46:52,765 INFO mapreduce.Job: The url to track the job: http://Azhar:8088/proxy/application_1696356523104_0001/
2023-10-03 23:46:52,770 INFO mapreduce.Job: Running job: job_1696356523104_0001
2023-10-03 23:47:07,174 INFO mapreduce.Job: Job job_1696356523104_0001 running in uber mode : false
2023-10-03 23:47:07,178 INFO mapreduce.Job: map 0% reduce 0%
2023-10-03 23:47:17,788 INFO mapreduce.Job: map 50% reduce 0%
2023-10-03 23:47:18,807 INFO mapreduce.Job: map 100% reduce 0%
2023-10-03 23:47:26,957 INFO mapreduce.Job: map 100% reduce 100%
2023-10-03 23:47:27,996 INFO mapreduce.Job: Job job_1696356523104_0001 completed successfully
```

STEP 16: At last Print your output for the MatrixMultiply text file.

Using the Command : hadoop fs -cat /outputmatrix/*

OUTPUT :

The screenshot shows a Windows Command Prompt window titled "Administrator: Command Prompt". The command entered is "hadoop fs -cat /outputmatrix/*". The output displays various performance metrics from the MapReduce framework, including time spent, vcore and megabyte usage, and memory statistics. It also lists shuffle errors and file format counters. The final output shows the matrix multiplication results:

```
Total time spent by all reduce tasks (ms)=3337
Total vcore-milliseconds taken by all map tasks=7652
Total vcore-milliseconds taken by all reduce tasks=3337
Total megabyte-milliseconds taken by all map tasks=7835648
Total megabyte-milliseconds taken by all reduce tasks=3417088
Map-Reduce Framework
  Map input records=8
  Map output records=8000
  Map output bytes=95120
  Map output materialized bytes=111132
  Input split bytes=202
  Combine input records=0
  Combine output records=0
  Reduce input groups=3996
  Reduce shuffle bytes=111132
  Reduce input records=8000
  Reduce output records=4
  Spilled Records=16000
  Shuffled Maps =2
  Failed Shuffles=0
  Merged Map outputs=2
  GC time elapsed (ms)=174
  CPU time spent (m)=795
  Physical memory (bytes) snapshot=1030872000
  Virtual memory (bytes) snapshot=1681747968
  Total committed heap usage (bytes)=906671488
  Peak Map Physical memory (bytes)=416120832
  Peak Map Virtual memory (bytes)=661319680
  Peak Reduce Physical memory (bytes)=280485888
  Peak Reduce Virtual memory (bytes)=486756352
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=68
File Output Format Counters
  Bytes Written=36
C:\hadoop-3.3.6\sbin>hadoop fs -cat /outputmatrix/*
0,0,19,0
0,1,22,0
1,0,43,0
1,1,50,0
C:\hadoop-3.3.6\sbin>
```

RESULT : Thus the program to run a basic wordcount mapreduce program to understand mapreduce is executed and output is verified successfully.